Vishnu Vardhan

Forward Deployed AI Engineer

7569449718 vishnudevv4518@gmail.com GitHub Trivandrum, India

Summary

Forward Deployed AI Engineer with **3 years of experience** working directly with clients to find the root cause behind vague requests (slow pipelines, manual workflows, unreliable AI outputs), then architecting and shipping the fix into live production with a measurable before/after result. Hands-on with **Agentic AI, RAG, vector databases (Pinecone/ChromaDB), LangChain, LangGraph, Python, FastAPI, REST APIs, SQL, Docker, AWS, and prompt engineering**. Delivered **hallucination-reduced RAG** and **multi-agent** systems to production, owning each project from diagnosis through deployed outcome.

Experience

Stealth Startup Bengaluru, Karnataka, India

Software Development Engineer 1 – Client-Facing AI & Backend Systems Aug 2023 – Sep 2026

- Partnered directly with clients across multiple live deployments to translate requirements for **document Q&A, semantic search, and automated reasoning** into **end-to-end AI backend pipelines** using **OpenAI/Gemini LLMs, LangChain agents, and RAG** on **Pinecone/ChromaDB**, shipped directly to **production**.
- Diagnosed slow, inconsistent backend APIs blocking client integrations and **architected 6+ production-grade REST APIs** with **Flask and FastAPI** using **modular MVC design**, cutting average response latency by **30%** across client environments.
- **Architected a multi-agent research and content pipeline** in **Python (LangChain, LangGraph)** with modular agent roles for web research, summarization, and drafting inside a stateful workflow, replacing manual, single-threaded content research and drafting.
- Built **autonomous multi-step reasoning** via **LangGraph state machines** with **tool-calling**, agent-to-agent delegation, and short/long-term memory, enabling agents to plan, execute, and self-correct without manual intervention.
- Integrated **OpenAI/Claude LLM APIs** with an **async FastAPI backend** for real-time triggering, monitoring, and retrieval of agent workflow results in production use.
- Reduced **LLM hallucination** by building a **RAG pipeline** on **Pinecone/ChromaDB** with custom embedding pipelines, grounding outputs in source data and improving factual accuracy.
- Resolved a latency bottleneck under multi-agent load using **Python asyncio orchestration, response caching, and parallelized agent calls**, reducing end-to-end pipeline latency.
- Root-caused slow query performance under load and **re-engineered MySQL schemas** (normalized design, composite indexing, stored procedures), cutting query execution time by **40%** under high-volume concurrent workloads.
- Closed a throughput ceiling during a client scale-up with **Apache Kafka event streaming, Redis caching, and Celery async task queues**, achieving **3× API throughput** and **sub-200ms p95** latency under **500+ concurrent requests**.
- **Re-architected a monolith into Dockerized microservices** with **Kafka messaging and Web3.py audit trails**, reducing deployment downtime by **50%** and enabling **zero-downtime rolling updates**.
- Owned security hardening ahead of a client compliance review (**JWT/OAuth2, RBAC, AES-256 encryption, parameterized queries**), delivering **zero critical CVEs** across all API-facing and blockchain endpoints.
- Served as the **cross-functional liaison** between clients and backend, AI/ML, blockchain, QA, and DevOps teams, sustaining **95%+ sprint velocity** across **8+ Agile cycles**.

Skills

FDE & Solution Delivery: Customer-Facing Engineering, Technical Consulting, Requirements Gathering, Business Requirements Analysis, Solution Architecture, System Integration, Production Deployment, Technical Troubleshooting, Root Cause Analysis, Performance Optimization, Stakeholder Management, Cross-Functional Collaboration
AI / LLM Engineering: Generative AI, LLM Integration (OpenAI GPT-4, Gemini, Claude, Groq), RAG, Prompt Engineering, Vector Embeddings, Semantic Search, Hybrid Search, Re-ranking, Hallucination Prevention, LLM Evaluation, AI Observability
Agentic AI: LangChain, LangGraph, MCP, Multi-Agent Systems, Agentic Workflows, Tool Calling, Agent-to-Agent Delegation, Autonomous Reasoning, Short/Long-Term Memory, Multi-LLM Routing & Fallback Strategies, Human-in-the-Loop
Backend & APIs: Python, Flask, FastAPI, REST/GraphQL API Design, Webhooks, Microservices, Asynchronous Programming (asyncio), Distributed Systems, Modular MVC
Data & Databases: MySQL, PostgreSQL (pgvector), MongoDB, Database Schema Design, Query Optimization, Pinecone, ChromaDB, Vector Databases, Database Integration
Cloud, DevOps & Infrastructure: AWS, AWS Bedrock, Docker, Kubernetes, CI/CD Pipelines, Cloud Deployment, Apache Kafka, Redis, Celery, Production Infrastructure, Monitoring, Logging, Observability
Enterprise Integration & Security: OAuth2, JWT, RBAC, Third-Party API Integration, Google Workspace, Slack, GitHub, Jira, Notion, Gmail API, Google Calendar API
Tools & Platforms: Git, GitHub, Postman, VS Code, n8n, Jupyter Notebook, Vercel, Netlify

Projects

Zero-Hallucination RAG Intelligence Platform | *React (Vite), Python, FastAPI, PostgreSQL (pgvector), OpenAI GPT-4, RAG, LLM*

- **Business problem:** standard LLM tools hallucinate on document-specific queries, producing confident but unverifiable answers with no traceability to source evidence.
- **Architecture:** architected and shipped a full-stack **Retrieval-Augmented Generation** platform integrating the **OpenAI GPT-4 API** with **hybrid dense + lexical search** over **PostgreSQL (pgvector)**, using **TanStack Start + React** for a server-rendered, production-grade frontend.
- Engineered a **6-layer hallucination-prevention system** — retrieval-confidence thresholds, clause-affinity re-ranking, evidence-overlap validation, prompt-contract enforcement, and post-generation **claim verification** — so the model **abstains rather than fabricates** when evidence is weak, missing, or conflicting.
- Built **multi-format ingestion** (PDF, DOCX, XLSX, CSV, TXT, Markdown, JSON, HTML) processing **421 indexed chunks across 3 live documents**, with automated chunking, embedding, and vector indexing delivering **traceable inline citations** on every generated claim.
- **Result:** benchmarked **0.0% hallucination rate** and **1.000 groundedness** across **40 evaluation queries** (24 answered, 16 correctly abstained) — with **1.000 recall@K**, **0.797 precision@K**, **1.000 MRR**, and **0.681 faithfulness**, at **2.3s average end-to-end latency**.
- **Next steps scoped:** deeper **database and third-party API integrations** to expand structured data grounding — pushing faithfulness and precision higher while extending **zero-hallucination** answers to live relational data. **Live App**

N45 Agentic Builder | — *Next.js, React, TypeScript, FastAPI, Python, MySQL, LLM, LangGraph, LangChain, RAG, Vector DB, AWS*

- **Business problem:** enterprises lacked a unified visual platform to build, orchestrate, and monitor multi-agent AI workflows — relying instead on fragmented scripts and manual integrations.
- **Architecture:** contributed to a production-ready **Enterprise AI Agent Development Platform** featuring a drag-and-drop **Workflow Builder (React Flow)** and **multi-agent orchestration** for parent–sub-agent collaboration.
- Worked on **workflow validation, versioning, autosave, deployment, and real-time execution monitoring** for reliable end-to-end workflow lifecycle management.
- Helped integrate **multiple LLM providers** (**OpenAI, Gemini, Anthropic, Groq, Ollama, Azure OpenAI**) via configurable AI profiles, with **Knowledge Base (RAG)** support for document ingestion and semantic retrieval.
- Supported implementation of **JWT authentication, RBAC, multi-tenant workspace architecture, and encrypted credential management** for enterprise-grade security.
- Assisted in connecting enterprise tools (**Google Workspace, Slack, GitHub, Jira, Notion, REST/GraphQL APIs, webhooks**) with **analytics dashboards and audit logging** for full observability.
- **Result:** platform enabled enterprises to build and scale AI agent workflows without custom infrastructure — cutting integration overhead across LLM providers and tools. **Live App**

Awards & Achievements

- Awarded **Best College Website Developer** by **Aditya Degree College** for excellence in website development .
- Collaborated with students from **NIT and IIT** to build and validate **business ideas** through cross-institution teamwork.
- Grew a social media audience to **20K+ followers**, building strong skills in content strategy, audience engagement, and digital outreach.

Education

Andhra University Visakhapatnam, Andhra Pradesh

Master of Computer Applications (MCA) – Pursuing